Vivekanand Education Society's

Institute of Technology

An Autonomous Institute Affiliated to University of Mumbai,, Approved by AICTE & Recognized by Govt. of Maharashtra

Hashu Advani Memorial Complex, Collector Colony, Chembur East, Mumbai - 400074.

Department of Information Technology

A.Y. 2024-25

Advance DevOps Lab
Experiment 12

Aim: To Build Your Application using AWS CodeBuild and Deploy on S3 / SEBS using AWS CodePipeline, deploy Sample Application on EC2 instance using AWS CodeDeploy.

Roll No.	44
Name	GANESH SANJAY PANDHRE
Class	D15B
Subject	Advance DevOps Lab
LO Mapped	LO1: To understand the fundamentals of Cloud Computing and be fully proficient with Cloud based DevOps solution deployment options to meet your business requirements. LO6: To engineer a composition of nano services using AWS Lambda and Step Functions with the Serverless Framework
Grade:	

AIM : To create a Lambda function which will log "An Image has been added" once you add an object to a specific bucket in S3.

THEORY :

AWS Lambda and S3 Integration

Delving into the powerful integration between AWS Lambda and Amazon S3, a popular use case in serverless architecture. AWS Lambda enables automatic execution of code in response to events generated by other AWS services, such as the addition of an object to an S3 bucket.

Event-Driven Architecture

AWS Lambda operates on an event-driven model, where specific events trigger the execution of predefined functions. In this case, when an object like an image is uploaded to a specific S3 bucket, S3 generates an event notification. This event is sent to AWS Lambda, which then triggers the function you have created. The function can be designed to perform various tasks, such as processing the object, extracting metadata, or, as in this experiment, simply logging a message—"An Image has been added."

Automation and Real-Time Processing

The integration of AWS Lambda with S3 is a prime example of automation in the cloud. It allows you to set up a real-time response system where actions are automatically triggered without manual intervention. This is particularly useful in scenarios where immediate processing or reaction is required, such as in applications that handle large volumes of incoming data or media files.

Scalability and Efficiency

One of the key advantages of using AWS Lambda with S3 is its scalability. Lambda functions can automatically scale to handle multiple events concurrently, ensuring that no matter how many objects are added to the S3 bucket, each will trigger the necessary function without delay. This setup is highly efficient, reducing operational overhead and ensuring that processes remain consistent and reliable, even under heavy workloads.

Use Cases and Practical Applications

This type of setup is widely used in various domains, including media processing, where images or videos uploaded to S3 can trigger Lambda functions to generate thumbnails or transcode files. It is also commonly used in data analytics, where data uploaded to S3 triggers processing functions that prepare the data for further analysis or storage. Another application is in backup systems, where new files added to S3 are immediately backed up or replicated in other locations.

Security Considerations

When working with AWS Lambda and S3, security is paramount. Each Lambda function interacts with S3 via permissions defined by AWS IAM roles and policies. These permissions must be carefully configured to ensure that the function has the necessary access rights—such as read and write permissions for the specific S3 bucket—without exposing unnecessary access to other resources. This ensures that the system remains secure while performing the required tasks efficiently.

Create an S3 Bucket. Enter a unique bucket name.

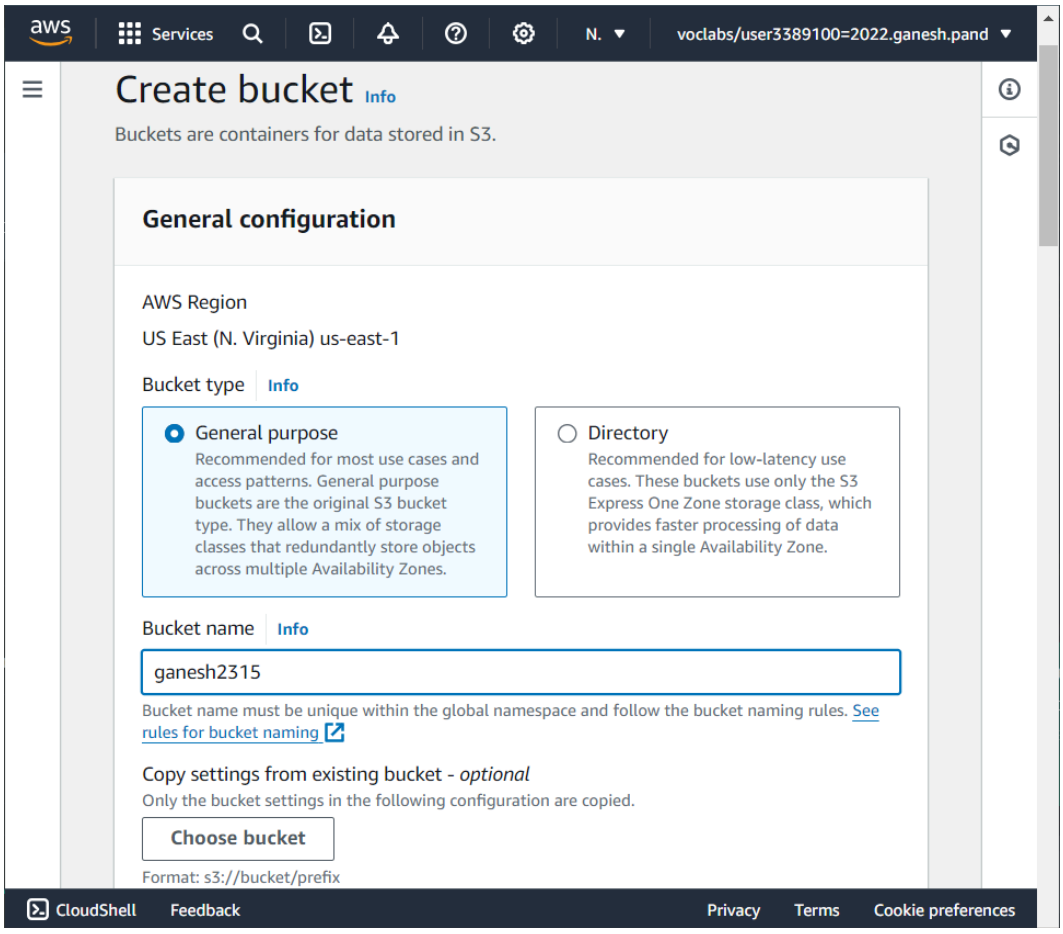

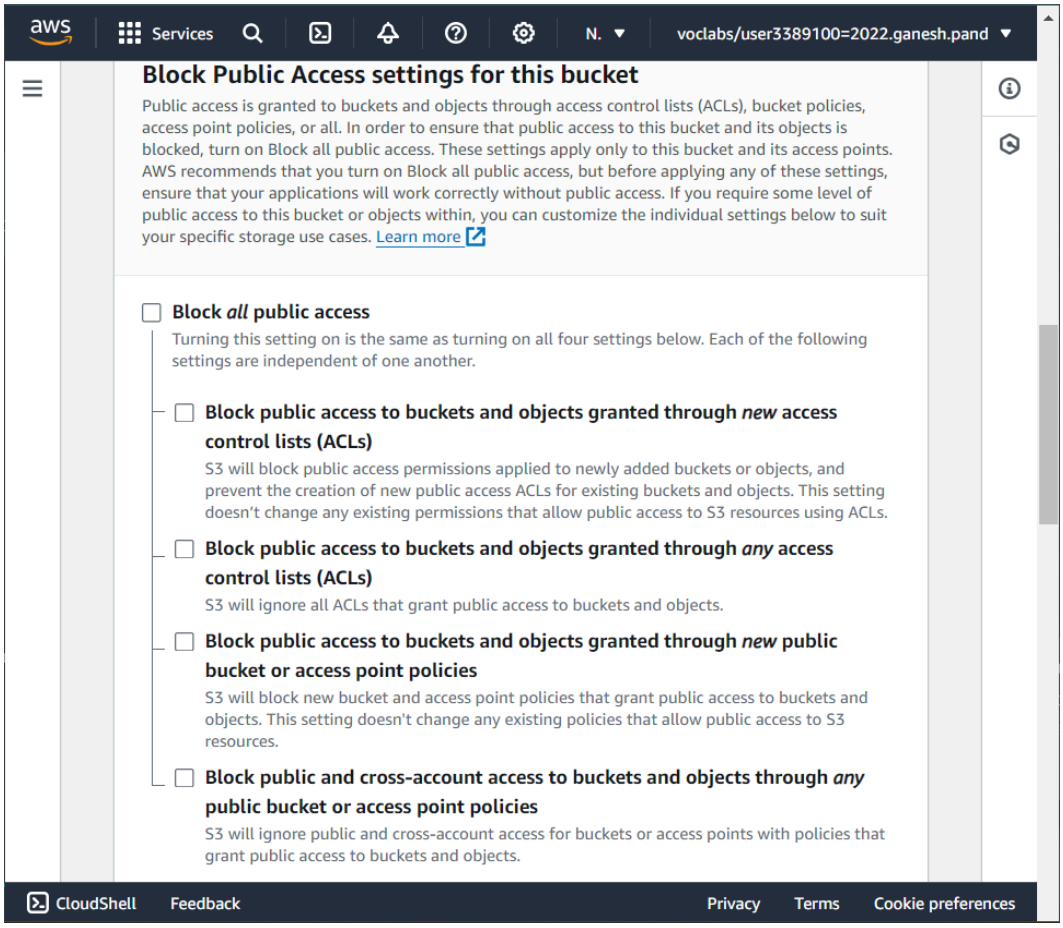

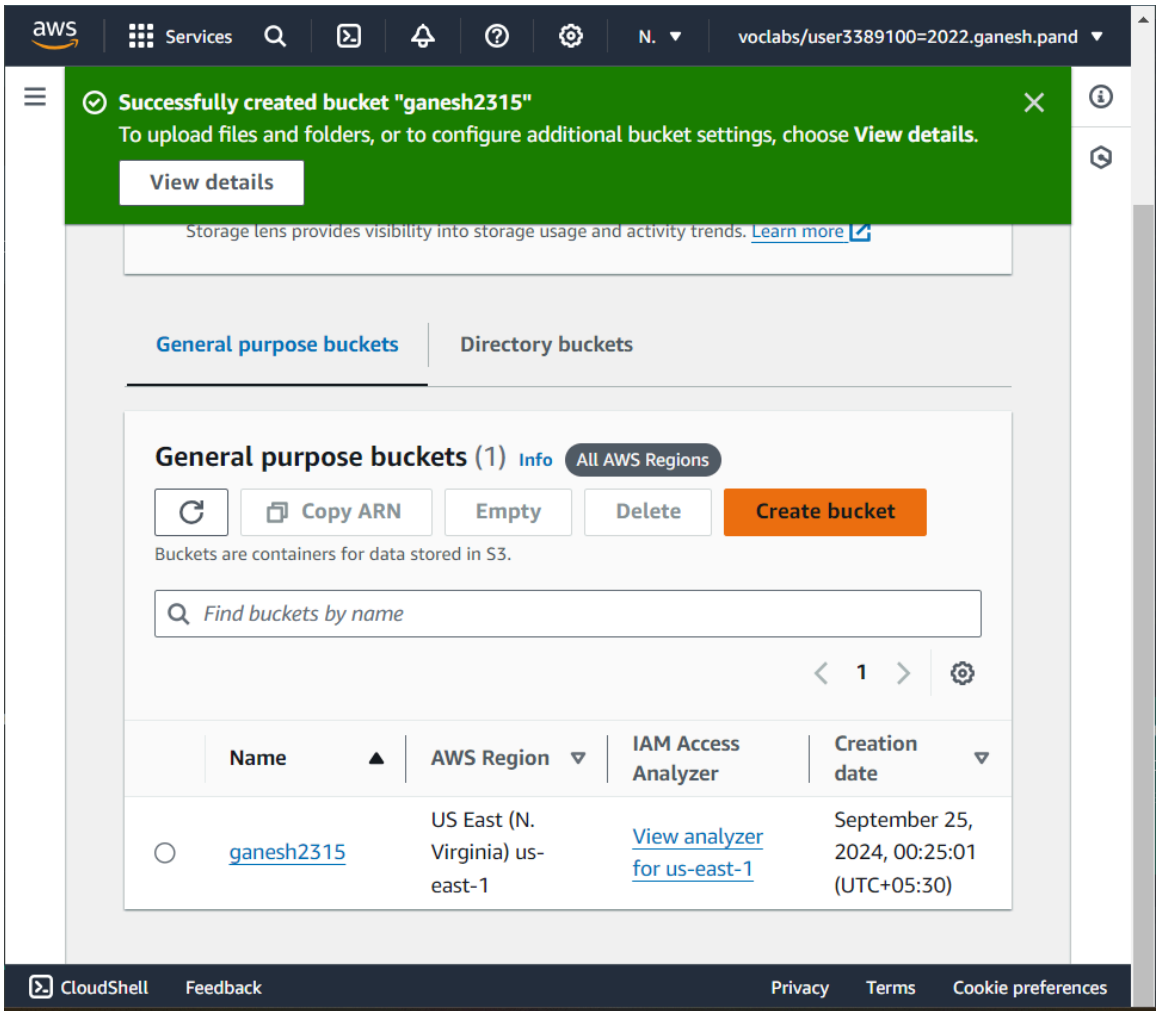

Navigate to the Lambda service. Click on "Create function."

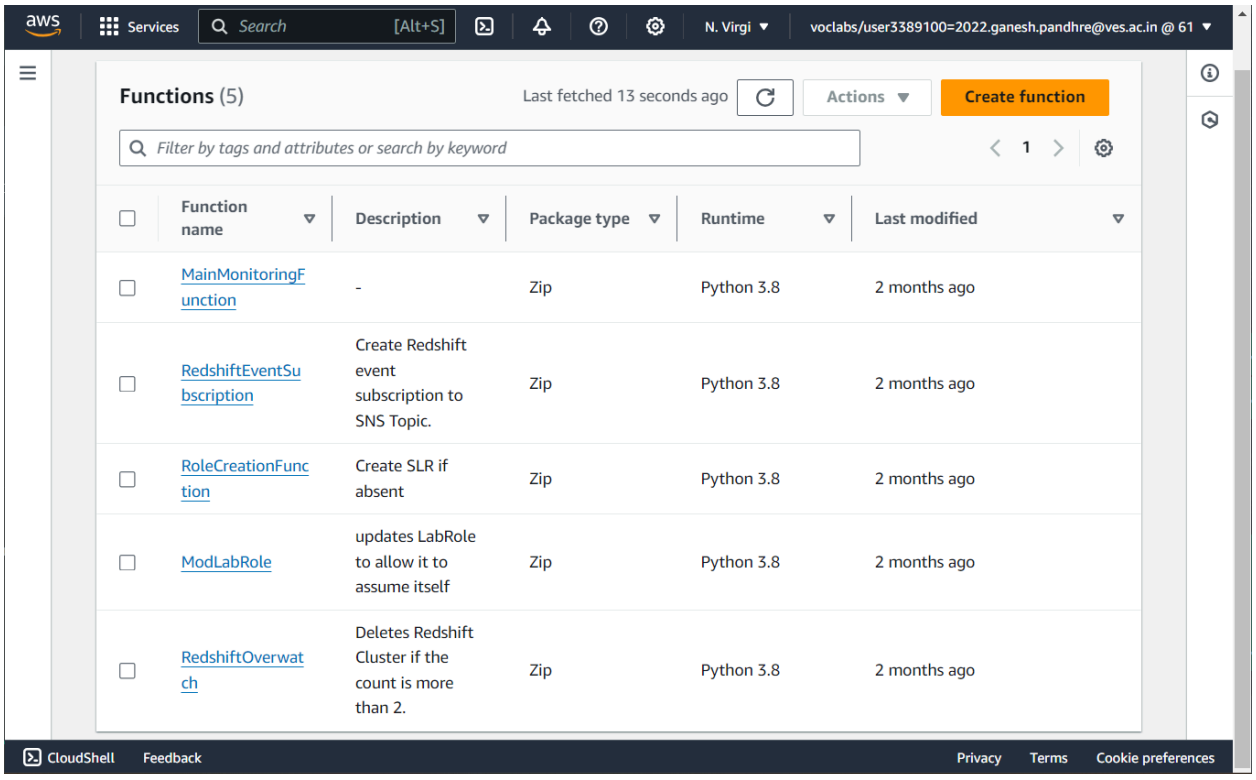

Choose "Author from scratch." Enter a name for your function. Select a runtime (Python 3.x).

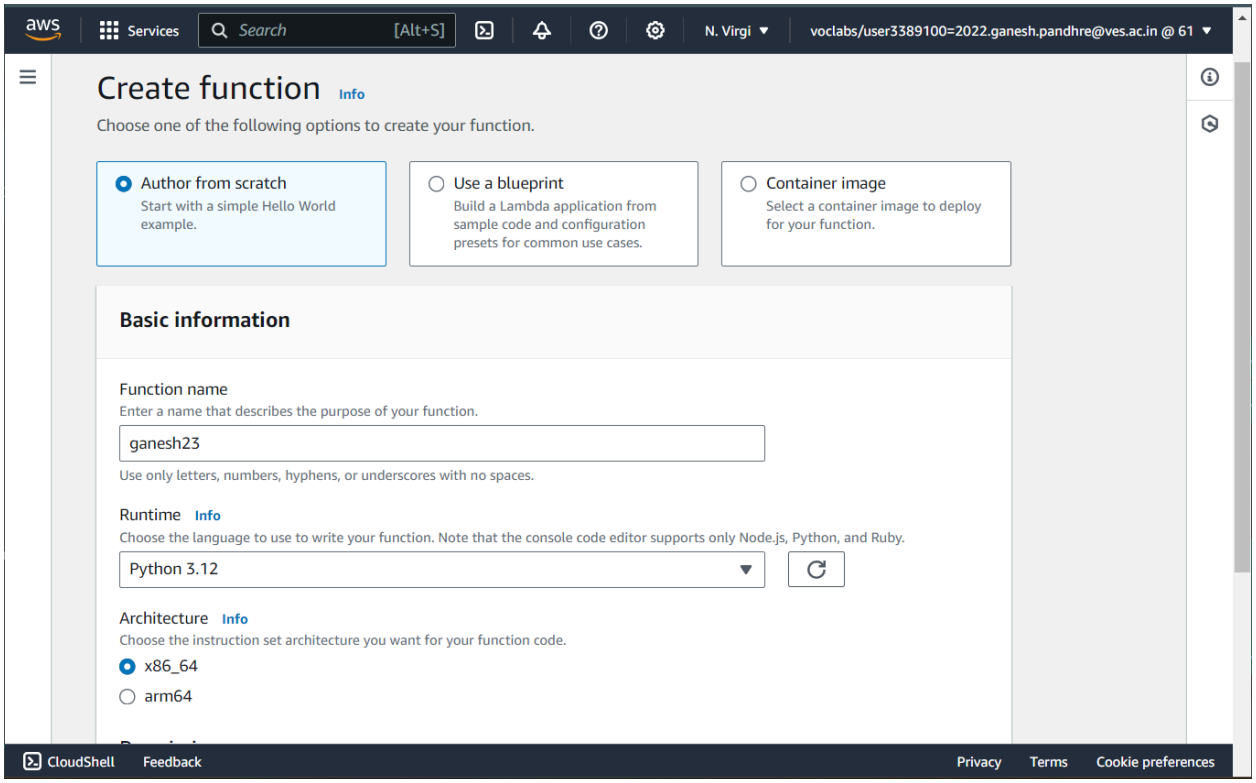

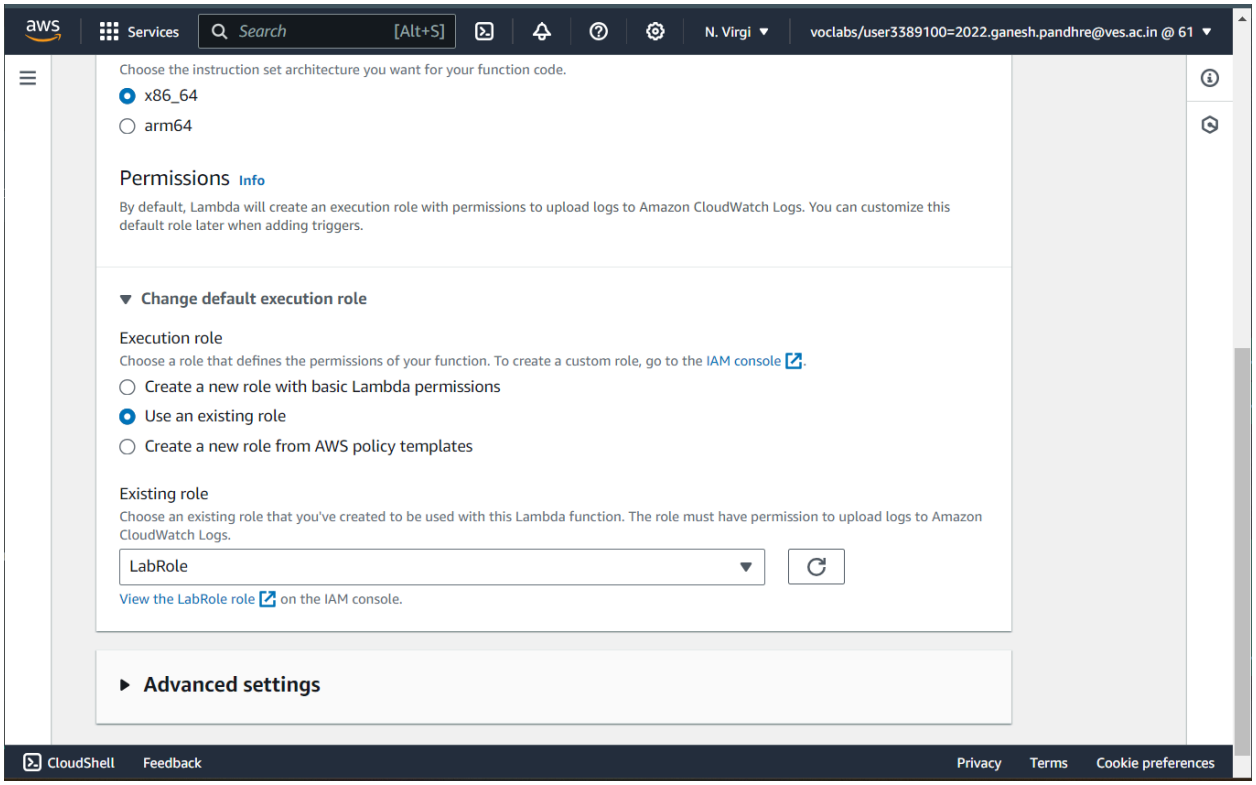

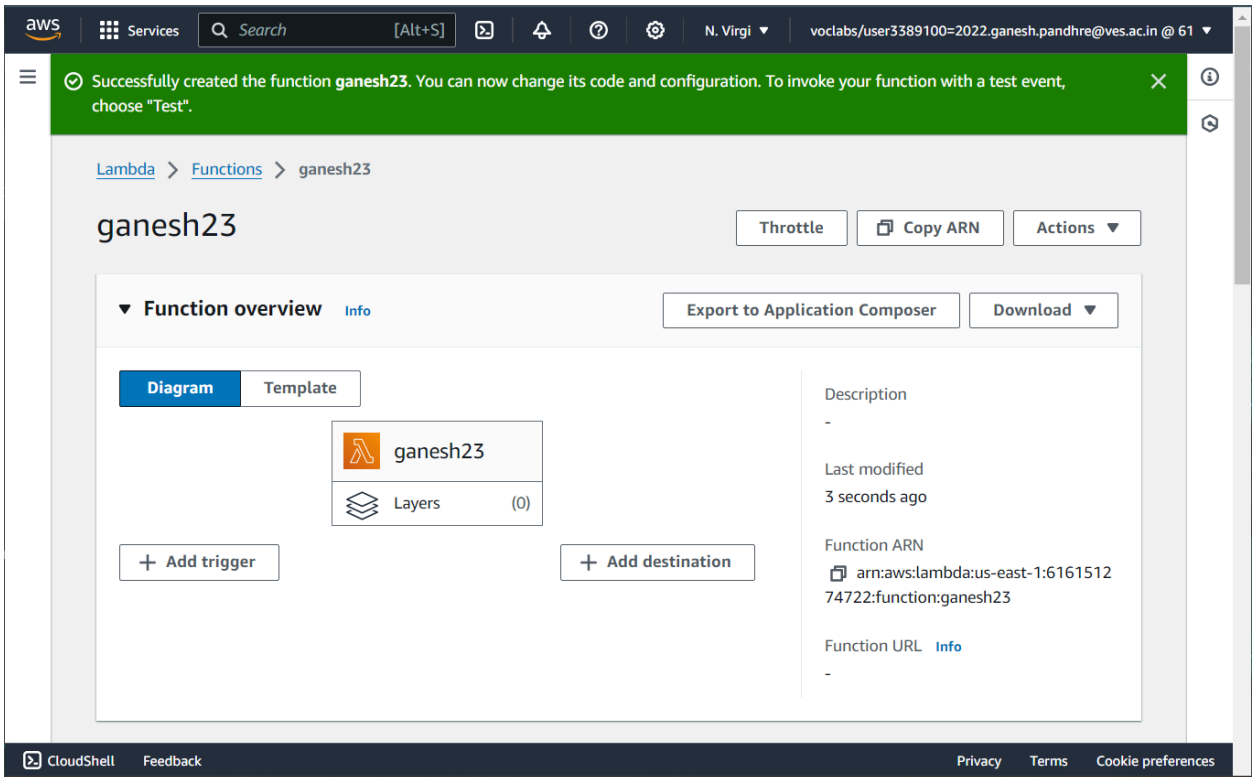

Write the Lambda Function Code. Click "Deploy" to save your changes.

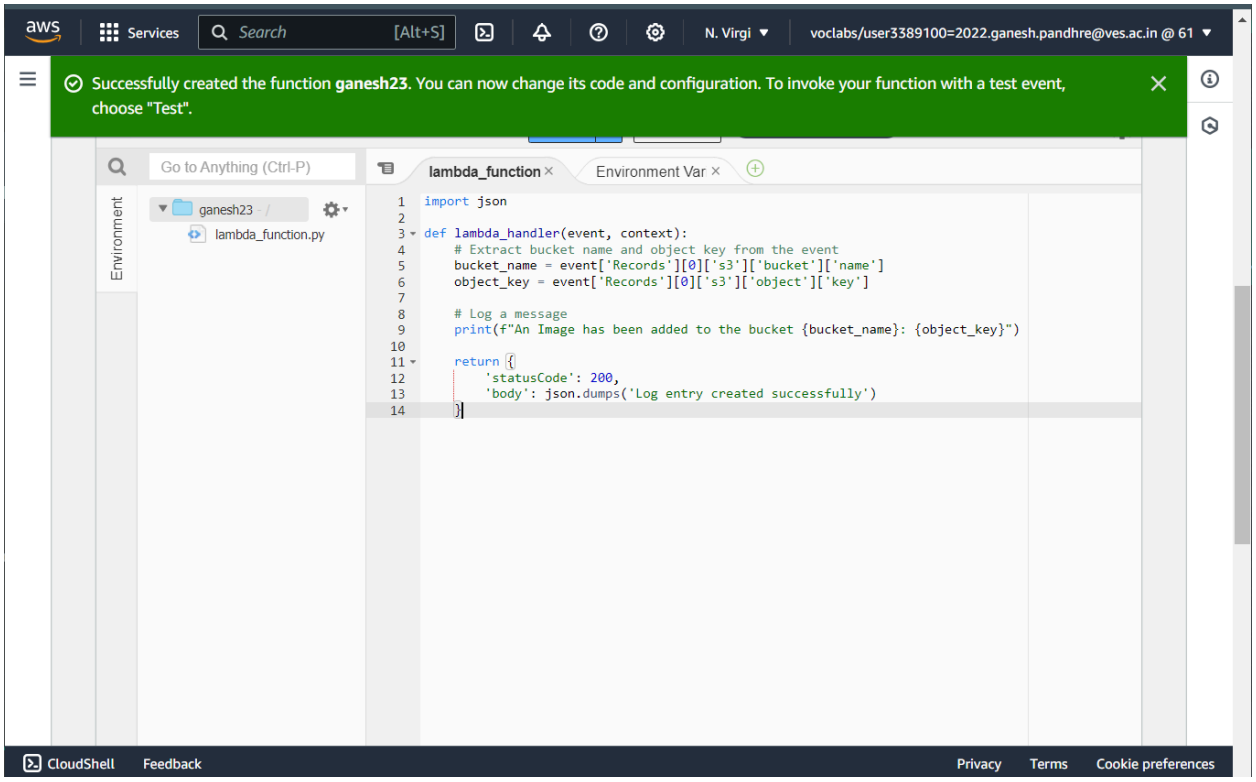

Select "S3" from the list of triggers. Choose the S3 bucket you created earlier. In the "Event type" dropdown, select "All object create events." Click "Add."

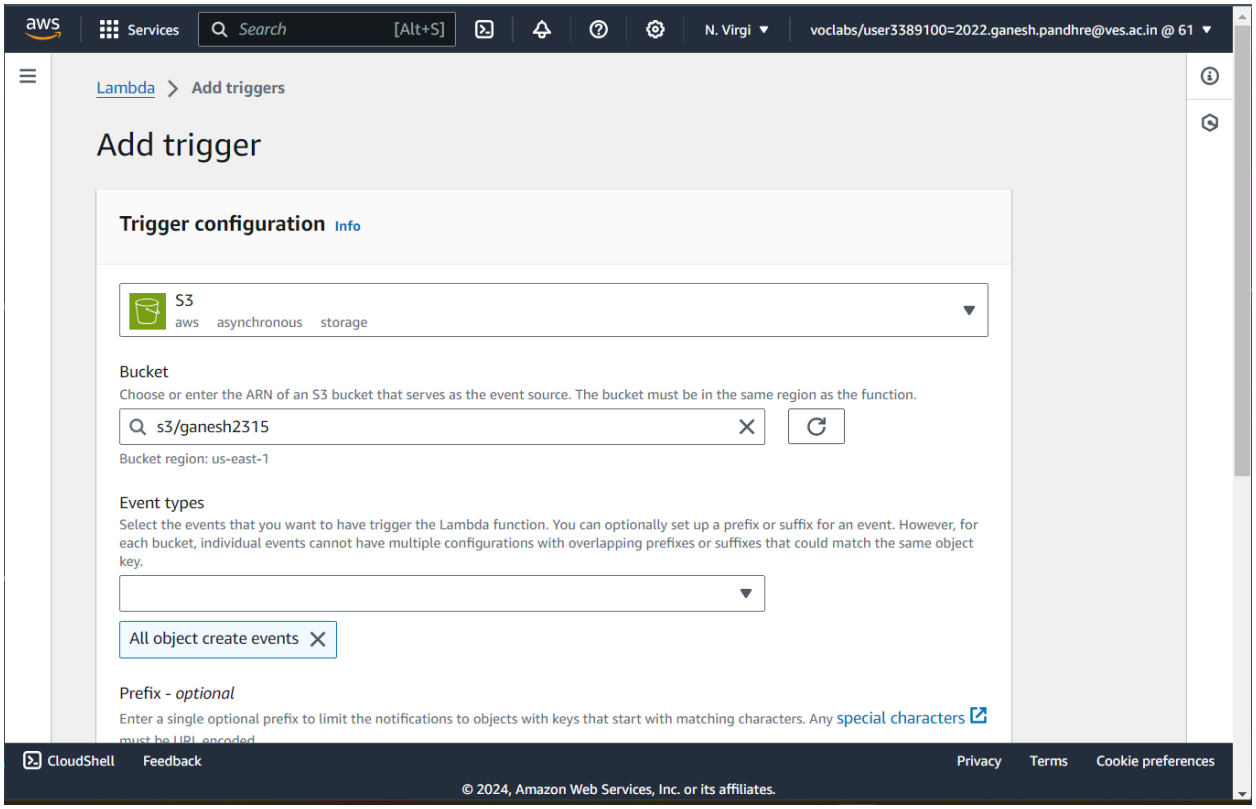

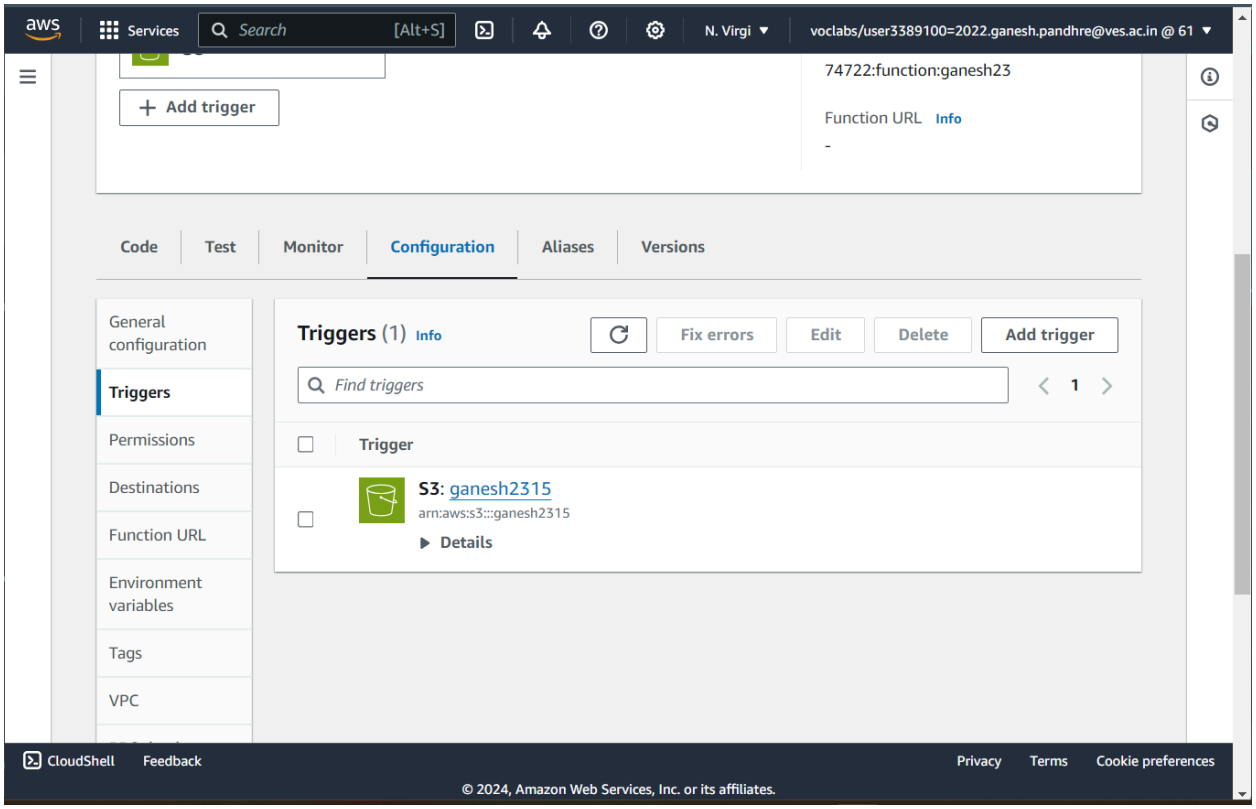

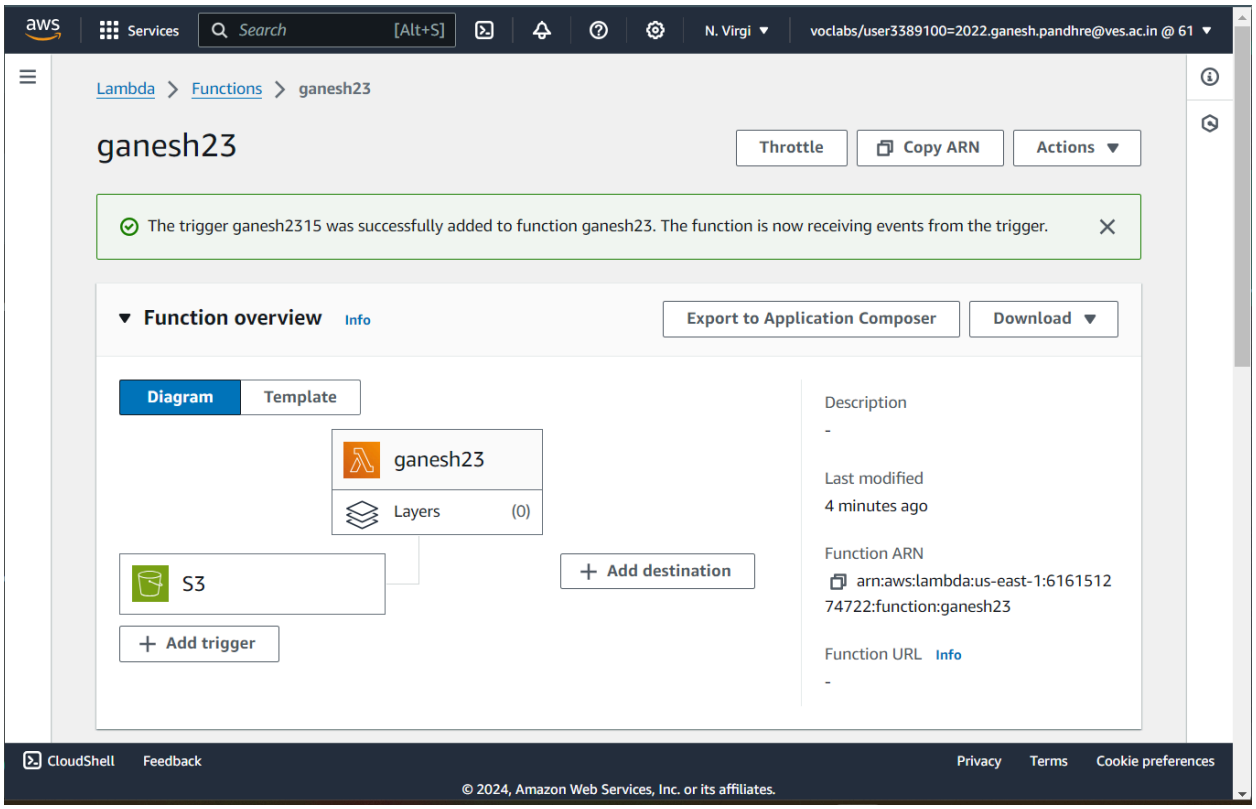

Navigate to the "Permissions" tab of your Lambda function. Ensure the Lambda function's execution role has the necessary permissions to access the S3 bucket.

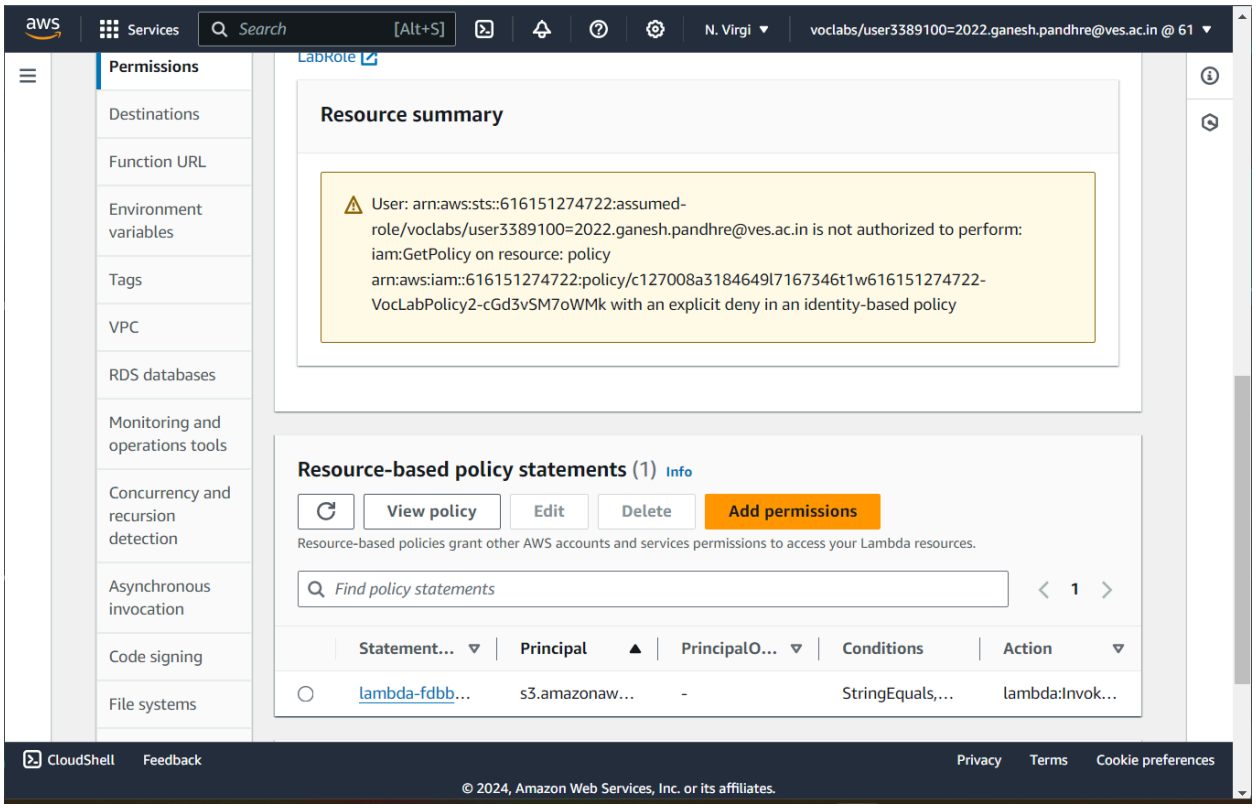

Upload an image file to your S3 bucket.

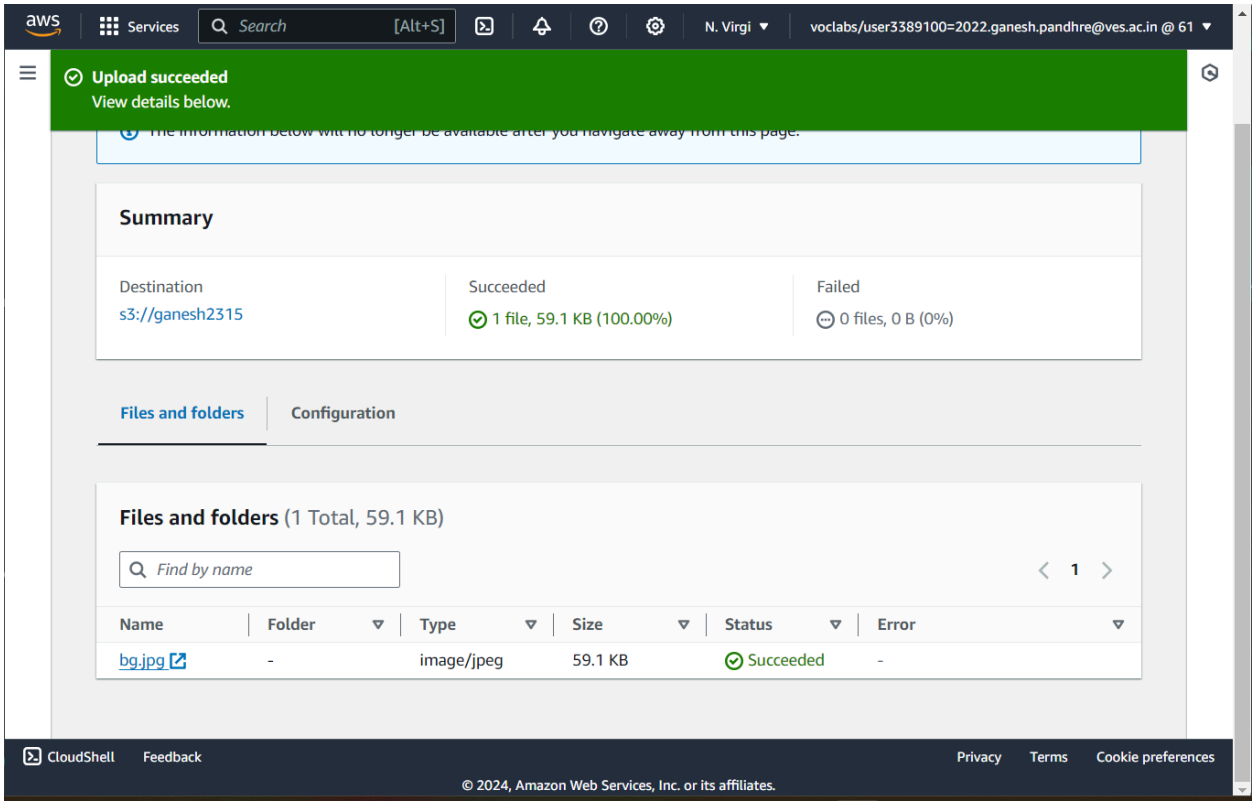

Go to the "Monitoring" tab in your Lambda function to check the logs. Use CloudWatch Logs to view the output and confirm that the message "An Image has been added" has been logged.

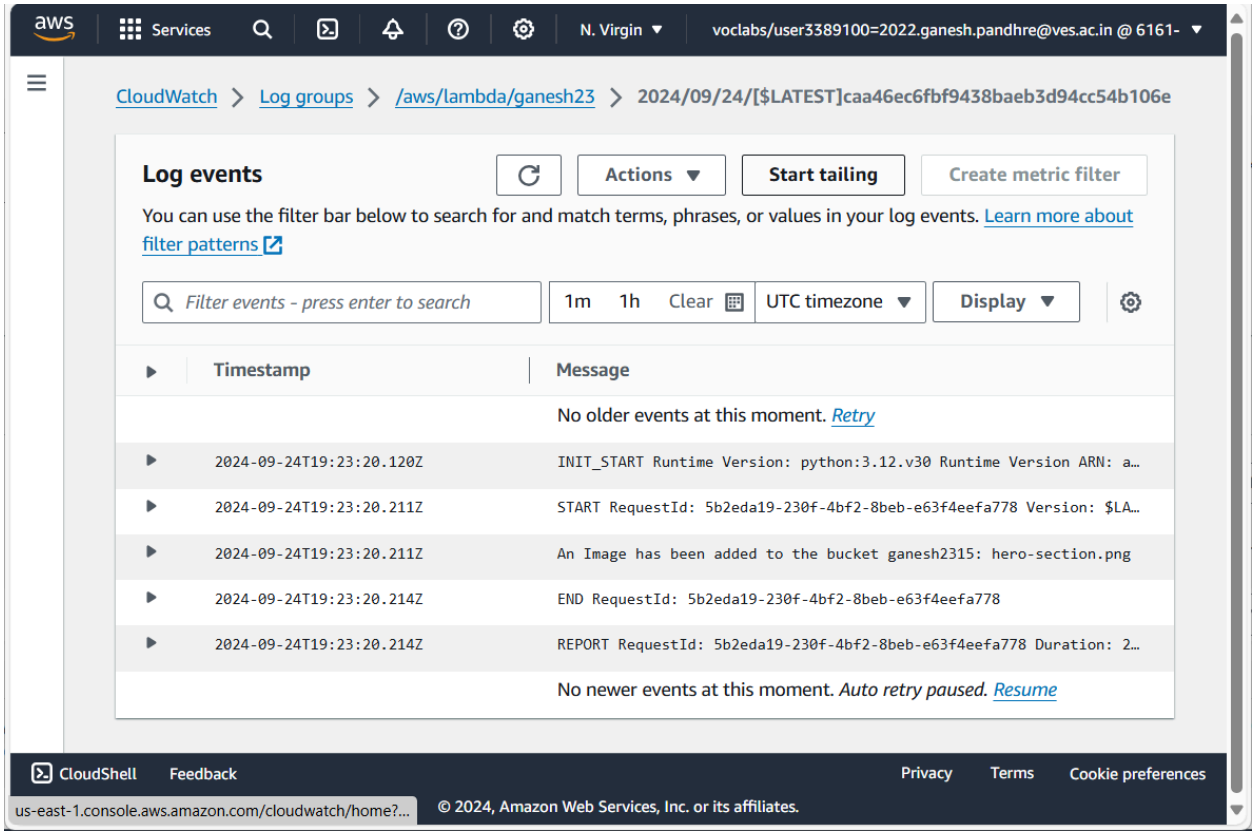

CONCLUSION : By integrating AWS Lambda with S3, you can create a seamless, automated process that responds instantly to changes in your S3 bucket. This not only improves efficiency but also allows for greater flexibility in handling data and media in real-time, making it a powerful tool in modern cloud-based applications.